

Transaction ID wraparound and avoiding the performance penalties from autovacuum tuple freezing

Martín Marqués

2ndQuadrant

2 de agosto de 2019

Contents

- 1 Introduction
- 2 Freezing
- 3 Autovacuum
- 4 Preventive vacuuming
- 5 Emergency shutdown
- 6 Conclusion

MVCC

- ▷ Tuple visibility is identified by create and delete txid
- ▷ Queries can't see uncommitted tuples
- ▷ Stricter isolation levels will only see data committed before the transaction started
- ▷ Hidden columns xmin and xmax, and tuple header
- ▷ Use modular 2^{31} arithmetic and tuple header to check visibility

Tuple identification

```
martin=> begin;
BEGIN
martin=> insert into test_id select s, clock_timestamp()
        from generate_series(1,10) s;
INSERT 0 10
martin=> select txid_current();
 txid_current
-----
          576
```

Tuple identification

```
martin=> select xmin,xmax, * from test_id ;
```

xmin	xmax	id	ts
576	0	1	2019-07-31 06:37:37.661925
576	0	2	2019-07-31 06:37:37.662361
576	0	3	2019-07-31 06:37:37.662384
576	0	4	2019-07-31 06:37:37.662391
576	0	5	2019-07-31 06:37:37.662397
576	0	6	2019-07-31 06:37:37.662404
576	0	7	2019-07-31 06:37:37.662409
576	0	8	2019-07-31 06:37:37.662415
576	0	9	2019-07-31 06:37:37.662421
576	0	10	2019-07-31 06:37:37.662428

Deleted tuples

Session 1

```
martin=> begin;
```

```
BEGIN
```

```
martin=> select txid_current();
```

```
txid_current
```

```
-----
```

```
577
```

```
martin=> update test_id set ts=now() where id>5;
```

```
UPDATE 5
```

Deleted tuples

Session 1

```
martin=> select xmin,xmax, * from test_id ;
```

xmin	xmax	id	ts
576	0	1	2019-07-31 06:37:37.661925
576	0	2	2019-07-31 06:37:37.662361
576	0	3	2019-07-31 06:37:37.662384
576	0	4	2019-07-31 06:37:37.662391
576	0	5	2019-07-31 06:37:37.662397
577	0	6	2019-07-31 06:39:57.201562
577	0	7	2019-07-31 06:39:57.201562
577	0	8	2019-07-31 06:39:57.201562
577	0	9	2019-07-31 06:39:57.201562
577	0	10	2019-07-31 06:39:57.201562

Deleted tuples

Session 2

```
martin=> select xmin, xmax, * from test_id ;
```

xmin	xmax	id	ts
576	0	1	2019-07-31 06:37:37.661925
576	0	2	2019-07-31 06:37:37.662361
576	0	3	2019-07-31 06:37:37.662384
576	0	4	2019-07-31 06:37:37.662391
576	0	5	2019-07-31 06:37:37.662397
576	577	6	2019-07-31 06:37:29.125214
576	577	7	2019-07-31 06:37:29.125214
576	577	8	2019-07-31 06:37:29.125214
576	577	9	2019-07-31 06:37:29.125214
576	577	10	2019-07-31 06:37:29.125214

Contents

- 1 Introduction
- 2 Freezing**
- 3 Autovacuum
- 4 Preventive vacuuming
- 5 Emergency shutdown
- 6 Conclusion

Which are the cons?

- ▷ txid are 32bit integers
- ▷ Only half of those txid are visible, that is 2147483648
- ▷ Highly transactional systems exhaust available txid

Which are the cons?

- ▷ txid are 32bit integers
- ▷ Only half of those txid are visible, that is 2147483648
- ▷ Highly transactional systems exhaust available txid
 - ▷ 100 tx/s → 8 months
 - ▷ 1000 tx/s → 24 days

How does postgres wrap-around

- ▷ We Freeze!
- ▷ Freeze old tuples whose xmin is older than all running backend xmin horizon
- ▷ Frozen tuples are always visible

How does postgres wrap-around

- ▷ Vacuum and autovacuum take care of freezing tuples

How does postgres wrap-around

- ▷ Vacuum and autovacuum take care of freezing tuples
 - ▷ `autovacuum_freeze_max_age`
 - ▷ `vacuum_freeze_table_age` and `vacuum_freeze_min_age`

How does postgres wrap-around

- ▷ Vacuum and autovacuum take care of freezing tuples
 - ▷ `autovacuum_freeze_max_age`
 - ▷ `vacuum_freeze_table_age` and `vacuum_freeze_min_age`
- ▷ `VACUUM FREEZE` is `vacuum` with `vacuum_freeze_table_age` and `vacuum_freeze_min_age` both set to zero

frozenxid

```
martin=> select relname, age(relfrozenxid),
           pg_size_pretty(pg_relation_size(oid)) as size
         from pg_class
         where relkind = 'r'
         order by age(relfrozenxid) desc, pg_relation_size(oid) desc
         limit 5;
```

relname	age	size
pg_proc	8029	608 kB
pg_depend	8029	448 kB
pg_collation	8029	432 kB
pg_attribute	8029	392 kB
pg_description	8029	320 kB

frozenxid

```

martin=> vacuum freeze pg_proc;
VACUUM
martin=> select relname, age(relfrozenxid),
               pg_size_pretty(pg_relation_size(oid)) as size
       from pg_class
       where relkind = 'r'
       order by age(relfrozenxid) desc, pg_relation_size(oid) desc
       limit 5;

```

relname	age	size
pg_depend	8029	448 kB
pg_collation	8029	432 kB
pg_attribute	8029	392 kB
pg_description	8029	320 kB
pg_operator	8029	120 kB

frozenxid

```
martin=> select relname, age(relfrozenxid),
              pg_size_pretty(pg_relation_size(oid)) as size
        from pg_class
        where relkind = 'r'
        order by age(relfrozenxid) asc, pg_relation_size(oid) desc
        limit 5;
```

relname	age	size
pg_proc	0	608 kB
test_id	0	224 kB
pg_statistic	0	136 kB
pg_depend	8029	448 kB
pg_collation	8029	432 kB

frozenxid

```
martin=> select datname, age(datfrozenxid)
          from pg_database
          where not datistemplate;
```

datname		age
postgres		8029
martin		8029

Freezing issues

- ▷ Only tuples visible to *all* backends can get frozen
- ▷ Long running transactions hold back freezing
- ▷ On PG 9.5 and older vacuum freeze would scan the whole table
 - ▷ PG 9.6 added *allfrozen* bit to the visibility map

Contents

- 1 Introduction
- 2 Freezing
- 3 Autovacuum**
- 4 Preventive vacuuming
- 5 Emergency shutdown
- 6 Conclusion

Autovacuum work

- ▷ Normal autovacuum will cancel itself if there are lock conflicts
- ▷ If `age(relfrozenxid) > autovacuum_freeze_max_age` autovacuum will run a **vacuum to prevent wraparound**
- ▷ An autovacuum to prevent wraparound will lock conflicting backends for as long as it's running
- ▷ Autovacuum to prevent wraparound will start at any time, very likely in your business peak hours
- ▷ Administrators increase `autovacuum_freeze_max_age` so autovacuum doesn't run these *annoying* vacuums. This just delays the inevitable

Autovacuum work

- ▷ On postgres 8.1 autovacuum didn't have freeze capabilities.

Autovacuum work

- ▷ On postgres 8.1 autovacuum didn't have freeze capabilities.
PLEASE UPGRADE!!!

Make Autovacuum Great (Again?)

- ▷ Default autovacuum parameters lax on aggressiveness
- ▷ Autovacuum needs to scan indexes as well

Make Autovacuum Great (Again?)

- ▷ Default autovacuum parameters lax on aggressiveness
- ▷ Autovacuum needs to scan indexes as well
- ▷ Big tables exacerbate the items above
- ▷ Too many indexes on a table will slow down autovacuum

Good news!

```
commit cbccac371c79d96c44fcd8c9cbb5ff4dedaaa522
```

```
Author: Tom Lane <tgl@sss.pgh.pa.us>
```

```
Date: Sun Mar 10 15:16:21 2019 -0400
```

Reduce the default value of `autovacuum_vacuum_cost_delay` to 2ms.

This is a better way to implement the desired change of increasing `autovacuum`'s default resource consumption.

Discussion: <https://postgr.es/m/28720.1552101086@sss.pgh.pa.us>

Even better good news!

```
commit a96c41feec6b6616eb9d5baee9a9e08c20533c38
```

```
Author: Robert Haas <rhaas@postgresql.org>
```

```
Date: Thu Apr 4 14:58:53 2019 -0400
```

Allow VACUUM to be run with index cleanup disabled.

This commit adds a new reloption, `vacuum_index_cleanup`, which controls whether index cleanup is performed for a particular relation by default. It also adds a new option to the VACUUM command, `INDEX_CLEANUP`, which can be used to override the reloption. If neither the reloption nor the VACUUM option is used, the default is true, as before.

Masahiko Sawada, reviewed and tested by Nathan Bossart, Alvaro Herrera, Kyotaro Horiguchi, Darafei Praliaskouski, and me.

The wording of the documentation is mostly due to me.

What can we do in the mean time?

- ▷ Make autovacuum more aggressive
- ▷ Partition large tables
- ▷ Remove unnecessary indexes
- ▷ Don't turn autovacuum off!!!

Contents

- 1 Introduction
- 2 Freezing
- 3 Autovacuum
- 4 Preventive vacuuming**
- 5 Emergency shutdown
- 6 Conclusion

What can we do in the mean time?

- ▷ Apply all the changes from the previous section on autovacuum
- ▷ During lower load hours run preventive vacuuming:
 - ▷ Query `pg_class` looking for tables with old `relfrozenxid`
 - ▷ `VACUUM` those tables with `vacuum_freeze_table_age` set to zero
 - ▷ Analyze the possibility of using more aggressive *cost* settings while vacuuming
- ▷ This prevents autovacuum from launching `VACUUM` to prevent wraparound during busier hours

Query for tables close to vacuum to prevent wraparound

```
SELECT c.oid::regclass as table,
       current_setting('autovacuum_freeze_max_age')::INT8 -
       age(c.relfrozenxid) as xid_left,
       pg_relation_size(c.oid) as relsize
FROM (pg_class c JOIN pg_namespace n ON (c.relnamespace=n.oid))
WHERE c.relkind = 'r' and
       age(c.relfrozenxid)::INT8 >
       (current_setting('autovacuum_freeze_max_age')::INT8 * 0.9)
ORDER BY 2 ASC
```

This query will only gather tables from a specific database so you will need to repeat the process for every database

Contents

- 1 Introduction
- 2 Freezing
- 3 Autovacuum
- 4 Preventive vacuuming
- 5 Emergency shutdown**
- 6 Conclusion

Once upon a time...

Once upon a time we consumed 2 billion txid's before autovacuum could freeze old tuples

- ▷ Server will send **WARNING** messages to the logs
- ▷ Once you reach wraparound system will effectively reject request for new txid
- ▷ Message will indicate shutting down, starting in single mode and running a database wide `VACUUM`

You are very close to wraparound territory

WARNING: database "martin" must be vacuumed within 123456789 transactions

HINT: To avoid a database shutdown, execute a database-wide VACUUM in that database.

You might also need to commit or roll back old prepared transactions, or drop stale replication slots.

You arrived to wraparound territory

ERROR: database is not accepting commands to avoid wraparound data loss in database "martin"

HINT: Stop the postmaster and vacuum that database in single-user mode.

You might also need to commit or roll back old prepared transactions, or drop stale replication slots.

How to get out of wraparound territory

```
/usr/pgsql-11/bin/postgres --single martin -D 11/data
```

```
PostgreSQL stand-alone backend 11.4
```

```
backend> VACUUM
```

```
backend>
```

Contents

- 1 Introduction
- 2 Freezing
- 3 Autovacuum
- 4 Preventive vacuuming
- 5 Emergency shutdown
- 6 Conclusion**

Conclusion

- ▷ Before postgres 8.2 you had to run manual `VACUUM` to freeze
- ▷ With 9.6 `VACUUM` can skip pages which are *allfrozen*
- ▷ Smaller tables are quicker to vacuum: Partition very big tables
- ▷ Postgres 12 adds feature to `INDEX_CLEANUP` when vacuuming
- ▷ `zheap` storage will likely eliminate all these problems when available in the future

Conclusion

- ▷ Before postgres 8.2 you had to run manual `VACUUM` to freeze
- ▷ With 9.6 `VACUUM` can skip pages which are *allfrozen*
- ▷ Smaller tables are quicker to vacuum: Partition very big tables
- ▷ Postgres 12 adds feature to `INDEX_CLEANUP` when vacuuming
- ▷ `zheap` storage will likely eliminate all these problems when available in the future

UPGRADE!!!

Conclusion

- ▷ Configure autovacuum with more aggressive *cost* values
- ▷ Run preventive vacuuming during lower load
- ▷ Monitor age of `relfrozenxid`
- ▷ Never set `autovacuum_freeze_max_age` to a value larger than 1 billion

Questions

?